

DOC-Lingo Test Document

Page 1 — Project Overview

Project Name: DOC-Lingo — Cross-Lingual Document Intelligence

Purpose: DOC-Lingo is a multilingual Retrieval-Augmented Generation (RAG) system designed to answer questions about documents. It allows a user to upload PDF or DOCX files and ask questions using English, Hindi, or Hinglish.

Core Capability: The system uses multilingual embeddings to represent documents and queries in a shared semantic space. This allows an English document to be searched using a Hindi or Hinglish question without requiring the original document to be translated before retrieval.

Technology Stack: Python, FastAPI, Next.js, React, LangChain, ChromaDB, Hugging Face multilingual embeddings, Ollama, and Retrieval-Augmented Generation.

How the System Works

First, the uploaded document is parsed and divided into meaningful chunks. Each chunk is converted into a multilingual vector embedding and stored in ChromaDB along with metadata such as the document name and page number.

When a user asks a question, the query is embedded using the same multilingual embedding model. ChromaDB searches for semantically relevant chunks. The retrieved context is then provided to the language model, which generates a grounded response.

Important test fact: DOC-Lingo should determine the answer language from the user's query or selected language, not from the language of the source document.

Expected Cross-Lingual Behavior

An English document should be searchable with an English query, a Hindi query, or a Hinglish query. The retrieved source content may remain in English, while the generated answer should match the requested response language.

DOC-Lingo Test Document

Page 2 — Organization and Operations

Organization: NovaTech Solutions is a fictional technology organization created specifically for testing DOC-Lingo. The company develops software for document management, data analysis, and intelligent information retrieval.

Founded: NovaTech Solutions was established in 2018 with its headquarters in Bengaluru. The organization initially focused on enterprise document processing and later expanded into artificial intelligence applications.

Employees: The company had 420 employees at the end of 2025. Its teams include software engineering, artificial intelligence research, product management, customer success, and information security.

2025 Business Information

During the 2025 financial year, NovaTech Solutions reported annual revenue of **■84 crore**. The company allocated approximately **■12 crore** to research and development. Its R&D; program focused on multilingual search, document understanding, and automated knowledge extraction.

The company served 185 enterprise customers across India, Singapore, and the United Arab Emirates. Its largest customer segment was financial services, followed by healthcare, education, and logistics.

DOC-Lingo Pilot Program

NovaTech introduced an internal DOC-Lingo pilot in October 2025. The objective was to help employees find information in policy documents, technical manuals, and internal reports using natural-language questions.

The pilot supported English documents and allowed employees to ask questions in English, Hindi, or Hinglish. The system returned answers together with document and page references so that users could verify the source information.

Key pilot result: Employees were able to retrieve relevant information from English-language documents using Hindi and Hinglish questions.

DOC-Lingo Test Document

Page 3 — Policies, Security, and Future Plans

Information Security Policy

NovaTech Solutions requires customer and internal documents to be handled according to its information security policy. Access to confidential documents is restricted to authorized employees, and document access is logged for auditing.

The DOC-Lingo pilot uses document-level filtering to ensure that users retrieve information only from documents available to their workspace. Sensitive credentials and API keys must never be stored inside uploaded documents or source code.

Future Development

The next phase of the project will focus on improving multilingual document support, better handling of tables, and more accurate retrieval from long documents. The team also plans to evaluate additional multilingual embedding models.

Another planned feature is improved language-aware generation. If a user asks a question in Hindi, the system should generate a Hindi response. If the user writes in Roman-script Hinglish, the system should produce a natural Hinglish response. An explicit language selector may override automatic language detection.

Suggested Questions for Testing

English: What was NovaTech Solutions' revenue in 2025?

Hindi: NovaTech Solutions ■■■ ■■■■■■■■ ■■ ■■■ ■■?

Hinglish: Company me 2025 ke end tak kitne employees the?

Cross-lingual: Is document ka DOC-Lingo pilot kab start hua tha?

Source test: Which page contains information about the DOC-Lingo pilot?

No-context test: What was NovaTech's 2026 revenue?

Expected Testing Outcome

The system should retrieve the correct page and answer using only the available document context. For Hindi and Hinglish questions, the final answer should remain in the requested language while preserving factual details such as dates, numbers, and company names.

